

[COM6513] Assignment 1: Sentiment Analysis with Logistic Regression

Instructor: Nikos Aletras

The goal of this assignment is to develop and test a **text classification** system for **sentiment analysis**, in particular to predict the sentiment of movie reviews, i.e. positive or negative (binary classification).

For that purpose, you will implement:

- Text processing methods for extracting Bag-Of-Word features, using
 - n-grams (BOW), i.e. unigrams, bigrams and trigrams to obtain vector representations of documents where $n=1,2,3$ respectively. Two vector weighting schemes should be tested: (1) raw frequencies (**1 mark**); (2) tf.idf (**1 mark**).
 - character n-grams (BOCN). A character n-gram is a contiguous sequence of characters given a word, e.g. for $n=2$, 'coffee' is split into {'co', 'of', 'ff', 'fe', 'ee'}. Two vector weighting schemes should be tested: (1) raw frequencies (**1 mark**); (2) tf.idf (**1 mark**). **Tip: Note the large vocabulary size!**
 - a combination of the two vector spaces (n-grams and character n-grams) choosing your best performing weighting respectively (i.e. raw or tfidf). (**1 mark**) **Tip: you should merge the two representations**
- Binary Logistic Regression (LR) classifiers that will be able to accurately classify movie reviews trained with:
 - (1) BOW-count (raw frequencies)
 - (2) BOW-tfidf (tf.idf weighted)
 - (3) BOCN-count
 - (4) BOCN-tfidf
 - (5) BOW+BOCN (best performing weighting; raw or tfidf)
- The Stochastic Gradient Descent (SGD) algorithm to estimate the parameters of your Logistic Regression models. Your SGD algorithm should:
 - Minimise the Binary Cross-entropy loss function (**1 mark**)
 - Use L2 regularisation (**1 mark**)
 - Perform multiple passes (epochs) over the training data (**1 mark**)
 - Randomise the order of training data after each pass (**1 mark**)
 - Stop training if the difference between the current and previous development loss is smaller than a threshold (**1 mark**)
 - After each epoch print the training and development loss (**1 mark**)
- Discuss how did you choose hyperparameters (e.g. learning rate and regularisation strength) for each LR model? You should use a table showing model performance using different set of hyperparameter values. (**2 marks**). ****Tip: Instead of using all possible combinations, you could perform a random sampling of combinations.**
- After training each LR model, plot the learning process (i.e. training and validation loss in each epoch) using a line plot. Does your model underfit, overfit or is it about right? Explain why. (**1 mark**).

- Identify and show the most important features (model interpretability) for each class (i.e. top-10 most positive and top-10 negative weights). Give the top 10 for each class and comment on whether they make sense (if they don't you might have a bug!). If you were to apply the classifier into a different domain such as laptop reviews or restaurant reviews, do you think these features would generalise well? Can you propose what features the classifier could pick up as important in the new domain? **(2 marks)**
- Provide well documented and commented code describing all of your choices. In general, you are free to make decisions about text processing (e.g. punctuation, numbers, vocabulary size) and hyperparameter values. We expect to see justifications and discussion for all of your choices **(2 marks)**.
- Provide efficient solutions by using Numpy arrays when possible (you can find tips in Lab 1 sheet). Executing the whole notebook with your code should not take more than 5 minutes on a any standard computer (e.g. Intel Core i5 CPU, 8 or 16GB RAM) excluding hyperparameter tuning runs **(2 marks)**.

Data

The data you will use are taken from here: <http://www.cs.cornell.edu/people/pabo/movie-review-data/> (<http://www.cs.cornell.edu/people/pabo/movie-review-data/>) and you can find it in the `./data_sentiment` folder in CSV format:

- `data_sentiment/train.csv`: contains 1,400 reviews, 700 positive (label: 1) and 700 negative (label: 0) to be used for training.
- `data_sentiment/dev.csv`: contains 200 reviews, 100 positive and 100 negative to be used for hyperparameter selection and monitoring the training process.
- `data_sentiment/test.csv`: contains 400 reviews, 200 positive and 200 negative to be used for testing.

Submission Instructions

You should submit a Jupyter Notebook file (assignment1.ipynb) and an exported PDF version (you can do it from Jupyter: `File->Download as->PDF via Latex` or you can print it as PDF using your browser).

You are advised to follow the code structure given in this notebook by completing all given functions. You can also write any auxiliary/helper functions (and arguments for the functions) that you might need but note that you can provide a full solution without any such functions. Similarly, you can just use only the packages imported below but you are free to use any functionality from the [Python Standard Library](https://docs.python.org/2/library/index.html) (<https://docs.python.org/2/library/index.html>), NumPy, SciPy (excluding built-in softmax functions) and Pandas. You are not allowed to use any third-party library such as Scikit-learn (apart from metric functions already provided), NLTK, Spacy, Keras etc..

There is no single correct answer on what your accuracy should be, but correct implementations usually achieve F1-scores around 80% or higher. The quality of the analysis of the results is as important as the accuracy itself.

This assignment will be marked out of 20. It is worth 20% of your final grade in the module.

The deadline for this assignment is **23:59 on Mon, 14 Mar 2022** and it needs to be submitted

via Blackboard. Standard departmental penalties for lateness will be applied. We use a range of strategies to **detect unfair means** (<https://www.sheffield.ac.uk/ssid/unfair-means/index>), including Turnitin which helps detect plagiarism. Use of unfair means would

```
In [1]: 1 import pandas as pd
2 import numpy as np
3 from collections import Counter
4 import re
5 import matplotlib.pyplot as plt
6 from sklearn.metrics import accuracy_score, precision_score, recall
7 import random
8 import string
9
10 # fixing random seed for reproducibility
11 random.seed(123)
12 np.random.seed(123)
13
14 # Generating 100 random words and their counts
```

Load Raw texts and labels into arrays

First, you need to load the training, development and test sets from their corresponding CSV files (tip: you can use Pandas dataframes).

```
In [2]: 1 # read csv file with the test sets
2 test = pd.read_csv("data_sentiment/test.csv")
3
4
5 # displaying the list of column names
6 #Column 1 = TEXT
7 #Column 2 = LABELS
8
9 # creating a list of column names by
10 # calling the columns
```

If you use Pandas you can see a sample of the data.

```
In [3]: 1 # read csv file with the train sets
2 train = pd.read_csv("data_sentiment/train.csv")
3
4
5 # displaying the list of column names
6 #Column 0 = TEXT
7 #Column 1 = LABELS
8
9
10 # creating a list of column names by
11 # calling the .columns
12 train_column_names = list(train.columns)
13
14
15 #####
16
17
18 # read csv file with the development sets
19 dev = pd.read_csv("data_sentiment/dev.csv")
20
```

```

21
22 # displaying the list of column names
23 #Column 0 = TEXT
24 #Column 1 = LABELS
25
26
27 # creating a list of column names by
28 # calling the .columns

```

The next step is to put the raw texts into Python lists and their corresponding labels into NumPy arrays:

```

In [4]: 1 #put the trainning raw texts into Python lists
2 train_text = list(train[train_column_names[0]])
3
4 #print the text for verification
5 #print(train_text,"\n")
6
7 #put the trainning labels into a NumPy arrays
8 train_label = train[train_column_names[1]].values
9
10 #print the train label for verification
11 print(train_label,"\n")
12
13
14 #####
15
16
17 #put the testing raw texts into Python lists
18 test_text = list(test[test_column_names[0]])
19
20 #print the text for verification
21 #print(test_text,"\n")
22
23 #put the testing labels into a NumPy arrays
24 test_label = test[test_column_names[1]].values
25
26 #print the test label for verification
27 print(test_label,"\n")
28
29
30 #####
31
32
33 #put the development raw texts into Python lists
34 dev_text = list(dev[dev_column_names[0]])
35
36 #print the text for verification
37 #print(dev_text,"\n")
38
39 #put the development labels into a NumPy arrays
40 dev_label = dev[dev_column_names[1]].values
41
42 #print the dev label for verification

```

[illegible]

To train and test Logistic Regression models, you first need to obtain vector representations for all documents given a vocabulary of features (unigrams, bigrams, trigrams).

To obtain a vocabulary of features, you should:

- tokenise all texts into a list of unigrams (tip: using a regular expression)
- remove stop words (using the one provided or one of your preference)
- compute bigrams, trigrams given the remaining unigrams (or character ngrams from the unigrams)
- remove ngrams appearing in less than K documents
- use the remaining to create a vocabulary of unigrams, bigrams and trigrams (or character n-grams). You can keep top N if you encounter memory issues.

```
1 #make a list for unigrams from each
2 def generate_ngrams(s, n):
3
4     # Convert to lowercases
5     #s = s.lower()
6
```

```

7      # Replace all none alphanumeric characters with spaces
8      s = re.sub(r'^a-zA-Z0-9\s', ' ', str(s))
9
10     # Break sentence in the token, remove empty tokens
11     tokens = [token for token in s.split(" ") if token != ""]
12
13     # Use the zip function to help generate n-grams
14     # Concatentate the tokens into ngrams and return
15     ngrams = zip(*[tokens[i:] for i in range(n)])
16     return [" ".join(ngram) for ngram in ngrams]
17
18     #Generate unigrams, using the trainning set
19     train_unigram = generate_ngrams(train_text, n=1)
20
21     #Generate unigrams, using the testing set
22     test_unigram = generate_ngrams(test_text, n=1)
23
24     #Generate unigrams, using the development set
25     dev_unigrams = generate_ngrams(dev_text, n=1)
26
27
28
29
30     stop_words = ['a', 'in', 'on', 'at', 'and', 'or',
31                  'to', 'the', 'of', 'an', 'by',
32                  'as', 'is', 'was', 'were', 'been', 'be',
33                  'are', 'for', 'this', 'that', 'these', 'those', 'you',
34                  'it', 'he', 'she', 'we', 'they', 'will', 'have', 'has',
35                  'do', 'did', 'can', 'could', 'who', 'which', 'what',
36                  'his', 'her', 'they', 'them', 'from', 'with', 'its']
37
38
39     #Remove these stop words from the list of ngrams
40     def remove_stop_word(ngram):
41
42         #use list comprehension,
43         #to only return words not inlcuded in stop_words
44         return [word for word in ngram if word not in stop_words]
45
46     #Remove stopwords on the train/test/dev unigrams
47     uni_train_no_sw = remove_stop_word(train_unigram)
48     uni_test_no_sw = remove_stop_word(test_unigram)
49     uni_dev_no_sw = remove_stop_word(dev_unigrams)
50
51     #Make Bigrams from the created unigrams
52     bi_train = generate_ngrams(uni_train_no_sw, n=2)
53     bi_test = generate_ngrams(uni_test_no_sw, n=2)
54     bi_dev = generate_ngrams(uni_dev_no_sw, n=2)
55
56     #Make Trigrams from the created unigrams
57     tri_train = generate_ngrams(uni_train_no_sw, n=3)
58     tri_test = generate_ngrams(uni_test_no_sw, n=3)
59     tri_dev = generate_ngrams(uni_dev_no_sw, n=3)
60
61
62     #####
63     #Remove ngrams appearing in less than K documents
64     def doc_counter(set_train, set_test, dev_test):
65
66         #use Counter type in order to count all unique words,

```

```

67     #from each train/test/dev set
68     c = Counter()
69     c.update(set_train)
70     c.update(set_test)
71     c.update(dev_test)
72
73     return c
74
75     #Initialise the set versions of the train/test/dev unigrams
76     set_uni_train = set(un_i_train_no_sw)
77     set_uni_test = set(un_i_test_no_sw)
78     set_uni_dev = set(un_i_dev_no_sw)
79
80     #Call doc_counter for all the unigram sets
81     uni_doc_appearances = doc_counter(set_uni_train, set_uni_test, set_uni_dev)
82
83     #Initialise the set versions of the train/test/dev bigrams
84     set_bi_train = set(bi_train)
85     set_bi_test = set(bi_test)
86     set_bi_dev = set(bi_dev)
87
88     #Call doc_counter for all the bigram sets
89     bi_doc_appearances = doc_counter(set_bi_train, set_bi_test, set_bi_dev)
90
91
92     set_tri_train = set(tri_train)
93     set_tri_test = set(tri_test)
94     set_tri_dev = set(tri_dev)
95
96     #Initialise the set versions of the train/test/dev trigrams
97     tri_doc_appearances = doc_counter(set_tri_train, set_tri_test, set_tri_dev)
98
99
100    #Number of documents(set between 1 and 3)
101    def find_words(c, k):
102
103        #Output list variable
104        found_words = []
105
106        #c is a Counter,
107        # go through every word contained by c
108        for words in c.keys():
109
110            #if documents appearance value is smaller than k
111            #in that case continue
112            if c[words] < k :
113                continue
114            else:
115                #Add this ngram to the list
116                found_words.append(words)
117
118            #No need to keep it as a list,
119            # arrays will help with efficiency
120            return np.array(found_words)
121
122
123    #Function that will return lists,
124    #containing only ngrams appearing at least in K documents
125    def remove_k(set_train, set_test, set_dev ,doc_ap):
126

```

```

127     clean_train = []
128     clean_test = []
129     clean_dev = []
130
131     print("Finding words....", "\n")
132
133     #Get all ngrams that appear in at least k documents
134     found_words = find_words(doc_ap, k=3)
135
136     print(found_words)
137
138     print("Starting trainning....", "\n")
139     clean_train = [word for word in set_train if word in found_words]
140
141     print("Starting testing....", "\n")
142     clean_test = [word for word in set_test if word in found_words]
143
144     print("Starting dev....", "\n")
145     clean_dev = [word for word in set_dev if word in found_words]
146
147     return np.array(clean_train), np.array(clean_test), np.array(clean_dev)
148
149 """

```

This code cell is separated from the rest due to the heavy computation needed.

Only run this cell once!

Estimated processing time with the full lists of ngrams: 12 minutes.

5 Minutes with the sets version

```

In [6]: 1
2
3
4     print("Starting UNIGRAMS....", "\n")
5
6
7     clean_uni_train, clean_uni_test, clean_uni_dev = remove_k(set_uni_train, set_uni_test, set_uni_dev, uni_doc_ap, k=3)
8
9
10
11
12
13     print("Starting BIGRAMS....", "\n")
14
15
16     clean_bi_train, clean_bi_test, clean_bi_dev = remove_k(set_bi_train, set_bi_test, set_bi_dev, bi_doc_ap, k=3)
17
18
19
20
21

```



```

22 print("Starting TRIGRAMS....", "\n")
23
24
25 clean_tri_train, clean_tri_test, clean_tri_dev = remove_k(set_tri_t
26                                                         set_tri_te
27                                                         set_tri_de
28                                                         tri_doc

```

Starting UNIGRAMS....

Finding words....

```
['emotive' 'macy' 'sketched' ... 'ms' 'virtues' 'aiming']
```

Starting trainning....

Starting BIGRAMS....

Finding words....

```
['no spark' 'take back' 'one films' ... 'every movie' 'friend amazin
g'
'known actors']
```

Starting trainning....

Starting TRIGRAMS....

Finding words....

```
['still doesn t' 'movie going experience' 'few far between' ...
'pretty good but' 'scale 0 4' 'best thing about']
```

Starting trainning....

```

In [7]: 1 #Create a vocabulary of unigrams, bigrams and trigrams
        2 vocab = set(clean_uni_train)
        3 vocab.update(clean_uni_test)
        4 vocab.update(clean_uni_dev)
        5
        6 vocab.update(clean_bi_train)
        7 vocab.update(clean_bi_test)
        8 vocab.update(clean_bi_dev)
        9
       10 vocab.update(clean_tri_train)
       11 vocab.update(clean_tri_test)
       12 vocab.update(clean_tri_dev)
       13
       14
       15
       16 print(vocab, "\n")
       17

```

```
{'emotive', 'macy', 'sketched', 'shrewd', 'director lasse', 'deep ri
sing', 'very believable', 'afraid', 'there more', 'goof', 'parallel
', 'civilization', 'physical appearance', 'ricci', 's possible', 'de
picting', 'covert', 'so ll', 'when', 'no sense all', 'no spark', 'po
lice station', 'store', 'years s', 'time just', 'gentleman', 's fell
ow', 'if s', 'taste', 'rd', 'gumption', 'bitchy', 'but comes', 'even
some', 'but overall', 'right hand', 'dr evil s', 'anytime', 'sexual
relationship', 'knocks', 'focuses', 'd probably', 'universal studios
', 'take back', 'doesn t get', 'founded', 'their friends', 'cloud',
'as great', 'expected', 'really satisfied', 'love film', 'really good' }
```

N-gram extraction from a document

You first need to implement the `extract_ngrams` function. It takes as input:

- `x_raw`: a string corresponding to the raw text of a document
- `ngram_range`: a tuple of two integers denoting the type of ngrams you want to extract, e.g. (1,2) denotes extracting unigrams and bigrams.
- `token_pattern`: a string to be used within a regular expression to extract all tokens. Note that data is already tokenised so you could opt for a simple white space tokenisation.
- `stop_words`: a list of stop words
- `vocab`: a given vocabulary. It should be used to extract specific features.
- `char_ngrams`: boolean. If true the function extracts character n-grams

and returns:

- `'x'`: a list of all extracted features.

See the examples below to see how this function should work.

```
In [8]: 1 def extract_ngrams(x_raw, ngram_range, token_pattern,
2         stop_words, vocab, char_ngrams):
3
4         #Set the smallest value of ngram types
5         min_ = ngram_range[0]
6
7         #Set the biggest value of ngram types
8         max_ = ngram_range[-1]
9
10        #Initialise output values
11        output_ngram = []
12        output_char_gram = []
13
14        #Produce Character ngrams or regular ngrams
15        if char_ngrams == False:
16
17            #Go through every type of ngram(i.e. unigram, bigram)
18            for rn in range(min_,max_+1):
19                print(rn)
20
21            # Replace all none alphanumeric characters with spaces
22            x_sub = re.sub(r'^a-zA-Z0-9\s]', ' ', str(x_raw))
23
24            x_sub.replace(" ", " ")
25
26            # Break sentence in the token, remove empty tokens
27            tokens = [token for token in x_sub.split(token_pattern)]
```

```

28
29         # Use the zip function to help generate n-grams
30         # Concatentate the tokens into ngrams and return
31         ngrams = zip(*[tokens[i:] for i in range(rn)])
32         final_ngrams = [" ".join(ngram) for ngram in ngrams]
33
34
35         #Remove stop words from ngrams
36         no_stop_ngram = [word for word in final_ngrams if word
37         #if rn == 3:
38             #print('This is the stop_words',*no_stop_ngram, sep=
39
40         #filter ngrams in vocabulary
41         for word_o in no_stop_ngram:
42             if word_o in vocab:
43                 output_ngram.append(word_o)
44
45         print(output_ngram)
46         return output_ngram
47
48     else:
49
50         #Generate character ngrams
51
52         #Go through every type of ngram(i.e. unigram, bigram)
53         for rn in range(min,max+1):
54
55             final_char =[]
56             #b[i:i+n] for i in range(len(b)-n+1)
57
58             # Replace all none alphanumeric characters with spaces
59             x_sub = re.sub(r"^[a-zA-Z0-9\s]", "", str(x_raw))
60
61             x_sub.replace("'", "")
62             x_sub.replace(" ", "")
63             # tokens = [token for token in x_sub.split(" ") if token
64
65
66             # Use the zip function to help generate character n-gra
67             # Concatentate the tokens into ngrams and return
68             char_grams = zip(*[x_sub[i:] for i in range(rn)])
69
70             #Split words by character, not by whitespace
71             final_char = ["".join(char_gram) for char_gram in char_
72
73
74             #Remove stopwords
75             output_char_gram = [word for word in final_char if word
76
77         print(output_char_gram)
78         return output_char_gram
79

```

Note that it is OK to represent n-grams using lists instead of tuples: e.g. ['great',
['great', 'movie']]

For extracting character n-grams the function should work as follows:

```
In [163]: 1 ##### Will Keep this cell commneted, for preview purposes #####
```

```

2
3 ##### To check the real running code go to cell above #####
4
5 # def extract_ngrams(x_raw="movie",
6 #                     ngram_range=(2,4),
7 #                     stop_words=[],
8 #                     char_ngrams=True):
9
10 #     min = ngram_range[0]
11
12 #     max = ngram_range[-1]
13
14 #     output_char_gram, no_stop_char = []
15
16
17 #     for rn in range(min,max+1):
18
19 #         #b[i:i+n] for i in range(len(b)-n+1)
20
21 #         # Replace all none alphanumeric characters with space
22 #         x_sub = re.sub(r'^a-zA-Z0-9\s|', ' ', str(x_raw))
23
24 #         # Break sentence in the token, remove empty tokens
25 #         tokens = [token for token in x_sub.split(token_pattern) if token]
26
27 #         # Use the zip function to help generate character n-grams
28 #         # Concatentate the tokens into ngrams and return
29 #         char_grams = zip(*[x_sub[i:] for i in range(rn)])
30 #         final_char = "".join(char_gram for char_gram in char_grams)
31
32
33 #         #Remove stopwords
34 #         no_stop_char = [word for word in final_char if word not in stop_words]
35
36
37 #         #search in vocab
38 #         output_char_gram = [word_o for word_o in no_stop_char if word_o in vocab]
39
40
41
42 #     return output_char_gram

```

Create a vocabulary

The `get_vocab` function will be used to (1) create a vocabulary of ngrams; (2) count the document frequencies of ngrams; (3) their raw frequency. It takes as input:

- `x_raw` : a list of strings each corresponding to the raw text of a document
- `ngram_range` : a tuple of two integers denoting the type of ngrams you want to extract, e.g. (1,2) denotes extracting unigrams and bigrams.
- `token_pattern` : a string to be used within a regular expression to extract all tokens. Note that data is already tokenised so you could opt for a simple white space tokenisation.
- `stop_words` : a list of stop words
- `min_df` : keep ngrams with a minimum document frequency.
- `keep_topN` : keep top-N more frequent ngrams.

and returns:

- vocab : a set of the n-grams that will be used as features.
- df : a Counter (or dict) that contains ngrams as keys and their corresponding document frequency as values.
- ngram_counts : counts of each ngram in vocab

Hint: it should make use of the `extract_ngrams` function

In [162]:

```

1  def get_vocab(X_raw, ngram_range, token_pattern,
2               min_df, keep_topN, stop_words, char_ngrams):
3
4      #Set the smallest value of ngram types
5      min_ = ngram_range[0]
6
7      #Set the biggest value of ngram types
8      max_ = ngram_range[-1]
9
10     ngrams = np.array([0])
11
12     #Go through every type of ngram(i.e. unigram, bigram)
13     for rn in range(min_, max_+1):
14         # n = ngram-range
15         print("Filter rn.... ", rn, "\n")
16
17
18         special_char=[" ", ":", " ", ";", ".", "?", "'"]
19
20         # Replace all none alphanumeric characters with spaces
21         s = re.sub(r'^a-zA-Z0-9\s', ' ', str(X_raw))
22
23         # Break sentence in the token, remove empty tokens
24         tokens = [token for token in s.split(" ") if token != ""]
25
26         # Use the zip function to help generate n-grams
27         # Concatentate the tokens into ngrams and return
28         n_grams = zip(*[tokens[i:] for i in range(rn)])
29         ngrams = np.append(ngrams, [" ".join(ngram) for ngram in n_grams])
30
31
32     print("Filter vocab....", "\n")
33
34     #Remove stop words and special charcters from the list of ngrams
35     filtered_vocab = [w for w in ngrams if w not in stop_words and
36                       char_ngrams(w)]
37
38
39     print("Start extract_ngram....", "\n")
40
41     #Initialise and pass the filtered vocab as a set
42     original_vocab = set(filtered_vocab)
43
44     #Extract ngrams from the vocabulary
45     ngram = extract_ngrams(X_raw, ngram_range, token_pattern, stop_words,
46                             char_ngrams)
47
48     #print(ngram)
49
50     #Count all the ngrams
51     df_count = Counter()
52     df_count.update(ngram)

```

```

52
53
54
55 def Compute_DF(ngrams):
56
57     DF = {}
58
59     print("Started DFs small.. \n")
60     for i in range(len(ngrams)):
61
62         for w in ngrams[i]:
63             try:
64                 DF[w].add(i)
65             except:
66                 DF[w] = {i}
67
68
69     for i in DF:
70         DF[i] = len(DF[i])
71     print(DF)
72     return DF
73
74
75 def find_doc_freq(word, DF):
76
77     #Method to get a specific ngram's Document frequency
78
79     c = 0
80
81     try:
82         c = DF[word]
83     except:
84         pass
85     return c
86
87
88
89 DF = Compute_DF(ngram)
90
91 vocab_init = set(ngram)
92
93 # found_gram = np.array([0])
94 #Filter ngrams through vocabulary PROBLEM
95 # found_gram = [w for w in ngram if w in vocab]
96
97
98 N = len(ngram)
99 # print(N)
100
101 #Initialise a new Counter to pass into ngrams with a count high
102 df_final = Counter()
103
104
105 df = {}
106
107
108 for i in range(N):
109
110     tokens = found_gram[i]
111     counter = Counter(tokens) #Replace with count vector

```

```

112         words_count = len(tokens)
113
114         #df_final.update(np.unique(tokens))
115
116         for token in np.unique(tokens):
117             tf = counter[token]/words_count
118             df_word = find_doc_freq(token, DF)
119             if df_word >= min_df:
120                 df.update({token: df_word}) #was df
121                 df_final.update(token)
122
123     vocab = set()
124     ngram_counts = []
125
126     #Go through the the top n most common ngrams,
127     # and extract their raw frequency and word
128
129     for word, count in df_final.most_common(keep_topN):
130         vocab.add(word)
131         ngram_counts.append(count) #Count is raw frequency
132
133
134     print(vocab)
135     print(df)
136     print(ngram_counts)
137     # print(type(top_ngrams))
138
139     # count = Counter()
140     # count.update(top_ngrams)
141
142
143     # ngram_counts = count.values()
144     # print(types(ngram_counts))
145

```

Now you should use `get_vocab` to create your vocabulary and get document and raw frequencies of n-grams:

```

In [164]: 1 test_vocab, test_df, test_count = get_vocab(test_text, ngram_range=
2           min_df=2, keep_topN=500,
3           stop_words = stop_words, char_ngrams = False)
4
5
6
7 # print('TEST VOCAB: ', test_vocab, '\n')
8

```

Filter rn.... 1

Filter rn.... 2

Filter rn.... 3

Filter vocab....

Start extract_ngram....

1
2
3

IOPub data rate exceeded.
The notebook server will temporarily stop sending output
to the client in order to avoid crashing it.
To change this limit, set the config variable
`--NotebookApp.iopub data rate limit`.

Then, you need to create 2 dictionaries: (1) vocabulary id -> word; and (2) word -> vocabulary id so you can use them for reference:

```
In [24]: 1 def create_2dict(df_dict):
2         id2word = {}
3         word2id = {}
4         dic_id = 0
5
6         for word in test_df.keys():
7
8             #(1) vocabulary id -> word
9             id2word.update({dic_id : word})
10
11             # (2) word -> vocabulary id
12             word2id.update({word: dic_id})
13
14             dic_id += 1
15
16         print('Dictionary [ID : WORD] : ',id2word, "\n")
17         print('Dictionary [WORD : ID] : ',word2id, "\n")
18
19         return id2word , word2id
```

Now you should be able to extract n-grams for each text in the training, development and test sets:

```
In [165]: 1 #TEST
2 # test_vocab, test_df, test_count = get_vocab(test_text, ngram_range=(1,3),
3 #                                             min_df=2, keep_topN=500,
4 #                                             stop_words = stop_words)
5
6 #train
7 train_vocab, train_df, train_count = get_vocab(train_text, ngram_range=(1,3),
8                                             min_df=10, keep_topN=100,
9                                             stop_words=stop_words,char_ngrams=False)
10 #Dev
11 dev_vocab, dev_df, dev_count = get_vocab(dev_text, ngram_range=(1,3),
12                                             min_df=10, keep_topN=100,
13                                             stop_words=stop_words,char_ngrams=False)
14
15
```



```
Filter rn.... 1
```

```
IOPubDataRateExceededError: Your IOPub data rate exceeded.
```

The notebook server will temporarily stop sending output

for the client in order to avoid crashing it.

To change this limit, set the config variable

```
c.NotebookApp.iopub_data_rate_limit`.
```

In [30]:

```
1 print("Test Dictionary >>> ", "\n")
2 id_w_test, w_id_test = create_2dict(test_df)
3
4 print("Train Dictionary >>> ", "\n")
5 id_w_train, w_id_train = create_2dict(train_df)
6
7 print("DEV Dictionary >>> ", "\n")
```

```
Test Dictionary >>>
```

```
Dictionary [ID : WORD] : {0: 'know', 1: 'but', 2: 'got', 3: 'around', 4: 'last', 5: 'one', 6: 'about', 7: 'final', 8: 'scene', 9: 'out', 10: 'enough', 11: 'watch', 12: 'such', 13: 'good', 14: 'behind', 15: 'show', 16: 'most', 17: 'well', 18: 'gets', 19: 'school', 20: 's', 21: 'plays', 22: 'him', 23: 'plot', 24: 'help', 25: 'very', 26: 'finds', 27: 'himself', 28: 'love', 29: 'fun', 30: 'begins', 31: 'go', 32: 'while', 33: 'goes', 34: 'like', 35: 'too', 36: 'young', 37: 't', 38: 'way', 39: 'two', 40: 'people', 41: 'really', 42: 'up', 43: 'year', 44: 'completely', 45: 'if', 46: 'me', 47: 'maybe', 48: 'best', 49: 'picture', 50: 'instead', 51: 'film', 52: 'reason', 53: 'so', 54: 'point', 55: 'next', 56: 'turn', 57: 'gives', 58: 'movie', 59: 'how', 60: 'having', 61: 'bit', 62: 'bad', 63: 'not', 64: 'being', 65: 'big', 66: 'woman', 67: 'performance', 68: 'had', 69: 'just', 70: 'because', 71: 'both', 72: 'into', 73: 'great', 74: 'character', 75: 'long', 76: 'when', 77: 'course', 78: 'actually', 79: 'even', 80: 'man', 81: 'every', 82: 'something', 83: 'life', 84: 'first', 85: 'friend', 86: 'where', 87: 'couple', 88: 'getting', 89: 'see', 90: 'anyt'}
```

Vectorise documents

Next, write a function `vectoriser` to obtain Bag-of-ngram representations for a list of documents. The function should take as input:

- `X_ngram`: a list of texts (documents), where each text is represented as list of n-grams in the `vocab`
- `vocab`: a set of n-grams to be used for representing the documents

and return:

- `X_vec`: an array with dimensionality $N \times |\text{vocab}|$ where N is the number of documents and $|\text{vocab}|$ is the size of the vocabulary. Each element of the array should represent the frequency of a given n-gram in a document.

In [152]:

```
1 def vectorise(X_ngram, vocab):
2
3
4     def count_vectorize(tokens):
```

```

5         ''' This function takes list of words in a sentence as inp
6         and returns a vector of size of filtered_vocab.It puts 0 i
7         word is not present in tokens and count of token if presen
8
9         vector = np.array([0])
10        for w in np.array(filtered_vocab):
11            vector = np.append(vector, tokens.count(w))
12
13
14        return vector
15
16    def Compute_DF(ngrams):
17
18        DF = {}
19
20        print("Started DFs small.. \n")
21        for i in range(len(ngrams)):
22
23            for w in ngrams[i]:
24                try:
25                    DF[w].add(i)
26                except:
27                    DF[w] = {i}
28
29
30        for i in DF:
31            DF[i] = len(DF[i])
32        print(DF)
33        return DF
34
35
36    def find_doc_freq(word,DF):
37
38        #Method to get a specific ngram's Document frequency
39
40        c = 0
41
42        try:
43            c = DF[word]
44        except:
45            pass
46        return c
47
48
49
50    def compute_tf_IDF(ngram):
51
52        #Calculate the Document Frequency
53
54        DF = Compute_DF(ngram)
55
56
57
58        doc = 0
59        token_counter = 0
60
61        print("Started TF.IDF...\n")
62
63        #Calculate TF.IDF
64

```

```

65
66     # N=Total number of documents in the dataset
67
68
69     found_gram = np.array([0])
70
71     #Filter ngrams through vocabulary PROBLEM
72     found_gram = [w for w in ngram if w in filtered_vocab]
73
74
75
76     print(found_gram)
77
78     N = len(found_gram)
79     print(N)
80
81     vocab_size = len(filtered_vocab)
82
83
84     dim_row = N
85     dim_columns = vocab_size
86
87
88     tf_idf = [[0 for j in range(dim_columns)] for i in range(d
89     print(tf_idf)
90
91     for i in range(N):
92
93         tokens = found_gram[i]
94         counter = Counter(tokens) #Replace with count vector
95         words_count = len(tokens)
96
97         token_counter = 0
98         for token in np.unique(tokens):
99
100
101             tf = counter[token]/words_count
102             df = find_doc_freq(token,DF)
103             idf = np.log(N/(df+1)) #numerator is added 1 to av
104
105             # df=total number of documents in which nth word o
106
107             tf_idf[i][token_counter] = tf*idf
108             token_counter +=1
109     #         tf_idf[doc, token] = tf*idf
110
111     #         doc += 1
112     #         token_counter += 1
113
114     print(np.array(tf_idf))
115     return np.array(tf_idf)
116
117
118
119
120
121
122
123     #list of special characters.You can use regular expressions to
124     special_char=[" ", ":", " ", " ", ";", ".", "?", "'"]

```

```

125
126
127     #split the sentences into tokens
128     x_sub = re.sub(r"^[a-zA-Z0-9\s]", " ", str(X_ngram))
129
130     tokens1 = [token for token in x_sub.split(" ") if token != ""]
131
132
133     #filter the vocabulary list
134     filtered_vocab = [w for w in vocab if w not in stop_words and
135
136     #print(filtered_vocab)
137
138     print("Count Vector...\n")
139     vector1=count_vectorize(tokens1)
140
141     print("Start compute_tf_IDF...\n")
142     TF_IDF_vector=compute_tf_IDF(tokens1)
143
144

```

Finally, use `vectorise` to obtain document vectors for each document in the train, development and test set. You should extract both count and tf.idf vectors respectively:

Count vectors

In [153]:

```

1  #COPY COUNT VECTORIZER HERE
2
3
4
5
6
7  #UNIGRAMS, SET_UNIGRAMS
8  print("Vectorise test text....","\n")
9  test_count, test_vect = vectorise(test_text, test_vocab)
10
11 print("Vectorise train text....","\n")
12 train_count, train_vect = vectorise(train_text, train_vocab)
13
14 print("Vectorise dev text....","\n")
15

```

Vectorise test text....

```
In [150]: 1 print('Array Shape = ',np.shape(test_vect) ) # test_vect.shape[0]
2 print('Array Shape = ',np.shape(train_vect) )
3
4 # print(np.shape(train_vect[0:19724, 0:]))
5
6 train_vect_sliced = train_vect[0:19724, 0:]#reshape
7 print('Array Shape = ',np.shape(train_vect_sliced) )
8
Array Shape = (72016, 500)
Array Shape = (141944, 100)
Array Shape = (19724, 100)
Array Shape = (19724, 100)
```

TF.IDF vectors

First compute `idfs` an array containing inverted document frequencies (Note: its elements should correspond to your `vocab`)

```
In [ ]: 1 #####COMPUTE TF.IDF using Term Frequency, Document Frequency and i
2
3 #Copy the TF.IDF method
4
5 #     tf_idf = {}
6 #     for i in range(N):
7 #         tokens = processed_text[i]
8 #         counter = Counter(tokens)
9 #         words_count = len(tokens)
10
11 #         for token in np.unique(tokens):
12 #             tf = count vector
13 #             df = doc_freq(token)
14
15 # -----> idf = np.log(N/(df+1)) <-----
16
17
18
19
20 # Formula can be one of these two:
21 #
22 # IDF = 1+log(N/dN)
23 #
24 # idf = log(N/(dN+1))
25
26 # Where
27
28 # N=Total number of documents in the dataset
```

Then transform your count vectors to tf.idf vectors:

```
In [ ]: 1 #         tf = counter[token]/words_count
2 # Replace with the count vector
3
In [ ]: 1 #         tf_idf[doc, token] = tf*idf
```

Binary Logistic Regression

After obtaining vector representations of the data, now you are ready to implement Binary Logistic Regression for classifying sentiment.

First, you need to implement the `sigmoid` function. It takes as input:

- `z` : a real number or an array of real numbers

and returns:

- `sig` : the sigmoid of `z`

```
In [34]: 1 def sigmoid(z):  
2  
3     sig = 1 / (1 + np.exp(-z))  
4     return sig  
5  
6     # z = np.dot(X, theta)  
7     # h = sigmoid(z)
```

Then, implement the `predict_proba` function to obtain prediction probabilities. It takes as input:

- `X` : an array of inputs, i.e. documents represented by bag-of-ngram vectors ($N \times |vocab|$)
- `weights` : a 1-D array of the model's weights ($1, |vocab|$)

and returns:

- `preds_proba` : the prediction probabilities of `X` given the weights

```
In [35]: 1 def predict_proba(X, weights):  
2  
3     preds_proba = sigmoid(np.dot(X, weights))  
4  
5  
6
```

Then, implement the `predict_class` function to obtain the most probable class for each vector in an array of input vectors. It takes as input:

- `X` : an array of documents represented by bag-of-ngram vectors ($N \times |vocab|$)
- `weights` : a 1-D array of the model's weights ($1, |vocab|$)

and returns:

- `preds_class` : the predicted class for each `x` in `X` given the weights

```
In [36]: 1 def predict_class(X, weights):  
2  
3
```

```

4
5     """
6 #         Predict the class between 0 and 1 using learned logistic
7 #         Using threshold value 0.5 to convert probability value to
8
9 #         I/P
10 #         -----
11 #         X : 2D array where each row represents a document and ea
12 #
13 #         weights : 1D array of weights. Dimension (1 x |vocab|)
14
15 #         O/P
16 #         -----
17 #         Class type based on threshold
18 #         """
19
20     p = preds_proba(X, weights) >= 0.5
21
22     preds_class = p.astype(int)
23
24 #     if y_pred_tr >= 0.5: #LABELS
25
26 #         predictions.append(1)
27 #     else:
28 #         predictions.append(0)
29
30     return preds_class

```

To learn the weights from data, we need to minimise the binary cross-entropy loss. Implement `binary_loss` that takes as input:

- `X` : input vectors
- `Y` : labels
- `weights` : model weights
- `alpha` : regularisation strength

and return:

- `l` : the loss score

```

In [37]: 1 def binary_loss(X, Y, weights, alpha=0.00001):
2
3     """
4 #         Compute cost for logistic regression.
5
6 #         I/P
7 #         -----
8 #         X : 2D array where each row represents a document and ea
9 #
10 #         y : 1D array of labels/target value for each training example
11 #
12 #         weights : 1D array of fitting parameters or weights. Dimension
13 #
14 #         alpha: regularisation strengths to be added when calculating
15 #
16 #         O/P
17 #         -----

```

```

18 #           J : The cost of using theta as the parameter for linear r
19 #           ""
20
21
22     m = len(X)
23     yhat = sigmoid(np.dot(X, weights) + alpha)
24
25     predict = Y * np.log(yhat) + (1 - Y) * np.log(1 - yhat)
26
27     l = -sum(predict) / m
28
29
30
31
32     return l
33
34
35

```

Now, you can implement Stochastic Gradient Descent to learn the weights of your sentiment classifier. The `SGD` function takes as input:

- `X_tr`: array of training data (vectors)
- `Y_tr`: labels of `X_tr`
- `X_dev`: array of development (i.e. validation) data (vectors)
- `Y_dev`: labels of `X_dev`
- `lr`: learning rate
- `alpha`: regularisation strength
- `epochs`: number of full passes over the training data
- `tolerance`: stop training if the difference between the current and previous validation loss is smaller than a threshold
- `print_progress`: flag for printing the training progress (train/validation loss)

and returns:

- `weights`: the weights learned
- `training_loss_history`: an array with the average losses of the whole training set after each epoch
- `validation_loss_history`: an array with the average losses of the whole development set after each epoch

```

In [174]: 1 def SGD(X_tr, Y_tr, X_dev, Y_dev, lr,
2           alpha, epochs,
3           tolerance, print_progress):
4
5     #           X = # data points with some features which we want to tr
6     #           y = # labels of all datapoints
7     #           # Initialize the weights and bias i.e. 'm' and 'c'
8     #           m = np.zeros_like(X[0]) # array with shape equal to no.
9     #           c = 0#regularisation
10    #           LR = 0.0001 # The learning Rate
11    #           epochs = 50 # no. of iterations for optimization
12
13
14    #           w=np.zeros(shape=(1,train_data.shape[1]-1))

```



```

15
16 #     C = f_integ(np.array([1]))
17 #     print "C", C
18 m_tr = np.zeros_like(X_tr)
19
20 m_dev = np.zeros_like(X_dev)
21
22 alpha_tr = alpha
23 alpha_dev = alpha
24
25
26 training_loss_history = np.array([0])
27 validation_loss_history = np.array([0])
28
29 training_loss_prev = np.array([0])
30 validation_loss_prev = np.array([0])
31
32 training_loss_current = np.array([0])
33 validation_loss_current = np.array([0])
34
35
36 # for every epoch
37 for epoch in range(1, epochs+1):
38
39     #####TRAINNING#####
40     # for every data point(X_train,y_train)
41     for i in range(len(X_tr)):
42
43         #compute gradient w.r.t 'm'
44         form_train = np.dot(X_tr[i], m_tr.T) + alpha_tr
45
46         gr_wrt_m_tr = X_tr[i]*(Y_tr[i] - sigmoid(form_train))
47
48         #compute gradient w.r.t 'c'
49         gr_wrt_c_tr = Y_tr[i] - sigmoid(form_train) #up
50
51         m_tr = m_tr - lr * gr_wrt_m_tr
52
53         alpha_tr = alpha_tr - lr * gr_wrt_c_tr# At the end of
54
55
56
57
58     if training_loss_prev == np.array([0]):
59
60         training_loss_prev = binary_loss(X_tr,Y_tr,m_tr,alpha_
61         training_loss_history = np.append(training_loss_histor
62
63     else:
64
65         training_loss_current = binary_loss(X_tr,Y_tr,m_tr,alp
66
67         if (training_loss_current - training_loss_prev) >= tol
68
69         training_loss_history = np.append(training_loss_hi
70         training_loss_prev = training_loss_current
71
72
73     #         if i % 10000 == 0:
74     if print_progress == True:

```

```

75         print("Loss after %d steps is: %.10f " % (epoch,traini
76
77     #####Development#####
78     # for every data point(X_train,y_train)
79     for j in range(len(X_dev)):
80
81         #compute gradient w.r.t 'm'
82         form_train = np.dot(X_dev[j], m_dev.T) + alpha_dev
83
84     #         In [1]: import numpy
85
86     #         In [2]: numpy.dot(numpy.ones([97, 2]), numpy.ones([2
87     #         Out[2]: (97, 1)
88
89         gr_wrt_m_dev = X_dev[j]*(Y_dev[j] - sigmoid(form_train
90
91         #compute gradient w.r.t 'c'
92         gr_wrt_c_dev = Y_tr[j] - sigmoid(form_train)          #u
93
94         m_dev = m_dev - lr * gr_wrt_m_dev
95
96         alpha_dev = alpha_dev - lr * gr_wrt_c_dev# At the end
97
98
99
100
101     if validation_loss_prev == np.array([0]):
102
103         validation_loss_prev = binary_loss(X_dev,Y_dev,m_dev,a
104         validation_loss_history = np.append(validation_loss_hi
105
106     else:
107
108         validation_loss_current = binary_loss(X_dev,Y_dev,m_de
109
110         if (validation_loss_current - validation_loss_prev) >=
111
112             validation_loss_history = np.append(validation_los
113             validation_loss_prev = validation_loss_current
114
115
116
117     #         validation_loss_history = np.append(validation_loss_hist
118
119     if print_progress == True:
120         print("Loss after %d steps is: %.10f " % (epoch,valida
121
122
123     #         weights
124
125     #         binary_loss(X_tr,Y_tr,m_tr,alpha_tr)
126
127     #         binary_loss(X_dev,Y_dev,m_dev,alpha_dev)
128
129     if print_progress == True:
130         print("Final loss after %d steps is: %.10f " % (epoch,traini
131         print("Loss after %d steps is: %.10f " % (epoch,validation
132         print("Final weights for training: ", m_tr,"\n")
133         print("Final weights for development: ", m_dev,"\n")
134

```

```

135 weights = np.array([0])
136 weights = np.append(weights, m_tr)
137 weights = np.append(weights, m_dev)
138
139 # So by using those optimum values of 'm' and 'c' we can perfo
140 #####MAYBE CALL predict class #####
141 #     for i in range(len(X_tr)):
142 #         z_tr = np.dot(X_tr[i], m) + alpha
143 #         y_pred_tr = sigmoid(z_tr)
144
145 #         if y_pred_tr>=0.5: #LABELS
146 #             predictions.append(1)
147 #         else:
148 #             predictions.append(0)
149
150 #     for i in range(len(X_dev)):
151
152 #         z_dev = np.dot(X_dev[i], m) + alpha
153 #         y_pred_dev = sigmoid(z_dev)
154
155 #         if y_pred_dev>=0.5:#LABELS
156 #             predictions.append(1)
157 #         else:
158 #             predictions.append(0)
159
160
161
162
163
164 # Make a prediction with coefficients
165 #     def predict(row, coefficients):
166 #         yhat = coefficients[0]
167 #         for i in range(len(row)-1):
168 #             yhat += coefficients[i + 1] * row[i]
169 #         return 1.0 / (1.0 + exp(-yhat))
170
171
172 #     # Estimate logistic regression coefficients using stochastic
173 #     def coefficients_sgd(train, l_rate, n_epoch):
174 #         coef = [0.0 for i in range(len(train[0]))]
175 #         for epoch in range(n_epoch):
176 #             for row in train:
177 #                 yhat = predict(row, coef)
178 #                 error = row[-1] - yhat
179 #                 coef[0] = coef[0] + l_rate * error * yhat * (1.0
180 #                 for i in range(len(row)-1):
181 #                     coef[i + 1] = coef[i + 1] + l_rate * error *
182 #             return coef
183
184 #     # Linear Regression Algorithm With Stochastic Gradient Desce
185 #     def logistic_regression(train, test, l_rate, n_epoch):
186 #         predictions = list()
187 #         coef = coefficients_sgd(train, l_rate, n_epoch)
188 #         for row in test:
189 #             yhat = predict(row, coef)
190 #             yhat = round(yhat)
191 #             predictions.append(yhat)
192 #         return(predictions)
193
194 #     def MyCustomSGD(train_data, learning_rate, n_iter, k, divideby):

```

```

195
196 #           # Initially we will keep our W and B as 0 as per the Tra
197 #           w=np.zeros(shape=(1,train_data.shape[1]-1))
198 #           b=0
199
200 #           cur_iter=1
201 #           while(cur_iter<=n_iter):
202
203 #               # We will create a small training data set of size K
204 #               temp=train_data.sample(k)
205
206 #               # We create our X and Y from the above temp dataset
207 #               y=np.array(temp['price'])
208 #               x=np.array(temp.drop('price',axis=1))
209
210 #               # We keep our initial gradients as 0
211 #               w_gradient=np.zeros(shape=(1,train_data.shape[1]-1))
212 #               b_gradient=0
213
214 #               for i in range(k): # Calculating gradients for point
215 #                   prediction=np.dot(w,x[i])+b
216 #                   w_gradient=w_gradient+(-2)*x[i]*(y[i]-(prediction))
217 #                   b_gradient=b_gradient+(-2)*(y[i]-(prediction))
218
219 #               #Updating the weights(W) and Bias(b) with the above
220 #               w=w-learning_rate*(w_gradient/k)
221 #               b=b-learning_rate*(b_gradient/k)
222
223 #               # Incrementing the iteration value
224 #               cur_iter=cur_iter+1
225
226 #               #Dividing the learning rate by the specified value
227 #               learning_rate=learning_rate/divideby
228
229 #           return w,b #Returning the weights and Bias
230 #####
231 #           class LogisticRegressionCustom():
232
233 #               def __init__(self, l_rate=1e-5, n_iterations=50000):
234 #                   self.l_rate = l_rate
235 #                   self.n_iterations = n_iterations
236
237 #               def initial_weights(self, X):
238 #                   self.weights = np.zeros(X.shape[1])
239
240 #               def sigmoid(self, s):
241 #                   return 1/(1+np.exp(-s))
242
243 #               #           m = len(X)
244 #               yhat = sigmoid(np.dot(X, weights) + alpha)
245
246 #               predict = Y * np.log(yhat) + (1 - Y) * np.log(1 - yhat)
247
248 #               l = -sum(predict) / m
249 #               return l
250
251 #               def binary_cross_entropy(self, X, y):
252 #                   return -(1/len(y)) * (y*np.log(self.sigmoid(np.dot(X,self.
253
254 #               def gradient(self, X, y):

```

```

255 #         return np.dot(X.T, (y-self.sigmoid(np.dot(X,self.weights
256
257 #     def fit(self, X, y):
258 #         self.initial_weights(X)
259 #         for i in range(self.n_iterations):
260 #             self.weights = self.weights+self.l_rate*self.gradien
261 #             if i % 10000 == 0:
262 #                 print("Loss after %d steps is: %.10f " % (i,self
263 #                 print("Final loss after %d steps is: %.10f " % (i,self.b
264 #                 print("Final weights: ", self.weights)
265 #         return self
266
267 #     def predict(self, X):
268 #         y_predict = []
269 #         for t in X:
270 #             y_predict.append(1) if self.sigmoid(np.dot(self.weig
271 #         return y_predict
272
273 #     def predict_proba(self, X):
274 #         y_predict = []
275 #         for t in X:
276 #             y_predict.append(self.sigmoid(np.dot(self.weights,t)
277 #         return y_predict
278
279 #####
280
281
282
283 #     def sigmoid(z):
284 #         sig = 1/(1+np.exp(-z))
285 #         return sig# Performing Gradient Descent Optimization
286
287
288
289
290
291
292     return weights, training_loss_history, validation_loss_history

```

Train and Evaluate Logistic Regression with Count vectors

First train the model using SGD:

In [175]:

```

1 print(type(train_vect))
2 print(np.shape(train_label))
3 print(type(train_count))
4
5 #BOW-count
6
7 weights, training_loss_history, validation_loss_history = SGD(train
8                                                         alpha
9                                                         toler
10
11 # (X_tr, Y_tr, X_dev, Y_dev, lr=0.1,
12 #     alpha=0.00001, epochs=5,
13 #     tolerance=0.0001, print_progress=True):
14
15 # print("Vectorise test text....","\n")
16 # test_vect = vectorise(test_text, test_vocab)

```

```

17 # print("Vectorise train text....","\n")
18 # train_vect = vectorise(train_text, train_vocab)
19 # print("Vectorise dev text....","\n")
20 # dev_vect = vectorise(dev_text, dev_vocab)
21
22
23 # print(train_label,"\n")
24
25
26 # #####
27
28
29 # #put the testing raw texts into Python lists
30 # test_text = list(test[test_column_names[0]])
31
32 # #print the text for verification
33 # #print(test_text,"\n")
34
35 # #put the testing labels into a NumPy arrays
36 # test_label = test[test_column_names[1]].values
37
38 # #print the label for verification
39 # print(test_label,"\n")
40
41
42 # #####
43
44
45 # #put the development raw texts into Python lists
46 # dev_text = list(dev[dev_column_names[0]])
47
48 # #print the text for verification
49 # #print(dev_text,"\n")
50
51 # #put the development labels into a NumPy arrays
52 # dev_label = dev[dev_column_names[1]].values
53
54 # #print the label for verification
55 # print(dev_label,"\n")

```

```
<class 'numpy.ndarray'>
```

```
(1399,)
```

```
<class 'list'>
```

```
<ipython-input-34-6527dd331435>:3: RuntimeWarning: overflow encountered in exp
```

```
    sig = 1 / (1 + np.exp(-z))
```

```
<ipython-input-37-8d7bf092e40c>:8: RuntimeWarning: divide by zero encountered in log
```

```
    predict = Y * np.log(yhat) + (1 - Y) * np.log(1 - yhat)
```

```
-----
```

```
ValueError                                Traceback (most recent call last)
```

```
<ipython-input-175-e176f2c34a9e> in <module>
```

```
5 #BOW-count
```

```
6
```

```
----> 7 weights, training_loss_history, validation_loss_history = SGD(train_count, train_label, dev_count, dev_label, lr=0.1,
```

```
8
```

```

alpha=0.00001, epochs=5,
9
tolerance=0.0001, print_progress=True)

<ipython-input-174-58f0580d6eb2> in SGD(X_tr, Y_tr, X_dev, Y_dev, l
r, alpha, epochs, tolerance, print_progress)
58     if training_loss_prev == np.array([0]):
59
---> 60         training_loss_prev = binary_loss(X_tr, Y_tr, m_tr,
alpha_tr)
61         training_loss_history = np.append(training_loss_
history, training_loss_prev)
62

<ipython-input-37-8d7bf092e40c> in binary_loss(X, Y, weights, alpha)
6     yhat = sigmoid(np.dot(X, weights) + alpha)
7
----> 8     predict = Y * np.log(yhat) + (1 - Y) * np.log(1 - yhat)
9
10     l = -sum(predict) / m

ValueError: operands could not be broadcast together with shapes (13
99,) (100,)

```

In []:

Now plot the training and validation history per epoch for the best hyperparameter combination. Does your model underfit, overfit or is it about right? Explain why.

In []:

```

1  # #plot
2
3  # from sklearn.metrics import roc_curve, roc_auc_score
4  # fpr, tpr, _ = roc_curve(y_test, y_prob)
5  # auc = roc_auc_score(y_test, y_prob)
6
7  # plt.figure(figsize=(10,8))
8  # plt.plot(fpr,tpr,label="data, auc="+str(round(auc,4)))
9
10 # plt.xlabel("False Positive Rate")
11 # plt.ylabel("True Positive Rate")
12
13 # plt.title("ROC Curve for Model from Scratch")
14 # plt.legend(loc=4)
15 # plt.show()
16
17 ##### MAIN #####
18 # training_loss_history
19 # validation_loss_history
20
21 plt.figure(figsize=(25,6))
22
23 plt.title('Cost Function Slope')
24 plt.plot(training_loss_history, label='Training Loss History')
25 plt.plot(validation_loss_history, label='Validation Loss History')
26 plt.legend(prop={'size': 16})
27 plt.xlabel('Number of Iterations')
28 plt.ylabel('Error Values')
29 plt.show()

```

```

30
31 ##### MAIN #####
32
33 # plt.figure(figsize=(10,8))
34 # plt.title('Cost Function Slope')
35 # plt.plot(cost)
36 # plt.xlabel('Number of Iterations')
37 # plt.ylabel('Error Values')

```

Explain here...

```

In [ ]: 1 #Underfit??
        2
        3 #Overfit??
        4
        5

```

Evaluation

Compute accuracy, precision, recall and F1-scores:

```

In [170]: 1 X_te_count = train_count
          2
          3 w_count = weights
          4
          5 preds_te_count = predict_class(X_te_count, w_count)
          6
          7 # train_count, weights
          8
          9
         10 Y_te = dev_count
         11
         12 print('Accuracy:', accuracy_score(Y_te,preds_te_count))
         13 print('Precision:', precision_score(Y_te,preds_te_count))
         14 print('Recall:', recall_score(Y_te,preds_te_count))
         15

```

```

-----
NameError                                Traceback (most recent call
1 last)
<ipython-input-170-f6a9ad36b6d1> in <module>
      1 X_te_count = train_count
      2
----> 3 w_count = weights
      4
      5 preds_te_count = predict_class(X_te_count, w_count)

NameError: name 'weights' is not defined

```

Finally, print the top-10 words for the negative and positive class respectively.

```

In [171]: 1 # id_w_test, w_id_test = create_2dict(test_df)
          2 # print("Train Dictionary >>> ", "\n")
          3
          4 # id_w_train, w_id_train = create_2dict(train_df)
          5 # print("DEV Dictionary >>> ", "\n")
          6

```



```

7 # id_w_dev, w_id_dev =create_2dict(dev_df)
8
9 top_neg = w_count.argsort()[ :10]
10 for i in top_neg:
11 #     print(id2word[i])
12

```

NameError Traceback (most recent call
l last)

```

<ipython-input-171-e06a79cbdd75> in <module>
      7 # id_w_dev, w_id_dev =create_2dict(dev_df)
      8
----> 9 top_neg = w_count.argsort()[ :10]
     10 for i in top_neg:
     11 #     print(id2word[i])

```

NameError: name 'w_count' is not defined

```

In [172]: 1 top_pos = w_count.argsort()[::-1][ :10]
          2 for i in top_pos:
          3 #     print(id2word[i])
          4

```

NameError Traceback (most recent call
l last)

```

<ipython-input-172-b8de8e434f3d> in <module>
----> 1 top_pos = w_count.argsort()[::-1][ :10]
      2 for i in top_pos:
      3 #     print(id2word[i])
      4     print(id_w_train)

```

NameError: name 'w_count' is not defined

If we were to apply the classifier we've learned into a different domain such laptop reviews or restaurant reviews, do you think these features would generalise well? Can you propose what features the classifier could pick up as important in the new domain?

Provide your answer here...

Sentiment Analysis

Discuss how did you choose model hyperparameters (e.g. learning rate and regularisation strength)? What is the relation between training epochs and learning rate? How the regularisation strength affects performance?

Enter your answer here...

(e.g. learning rate and regularisation strength)

Train and Evaluate Logistic Regression with TF.IDF

vectors

Follow the same steps as above (i.e. evaluating count n-gram representations).

Now repeat the training and evaluation process for BOW-tfidf, BOCN-count, BOCN-tfidf, BOW+BOCN including hyperparameter tuning for each model...

BOW-tfidf:

```
In [ ]: 1 #####BOW-tfidf#####
2
3 # #TEST
4 # test_vocab, test_df, test_count = get_vocab(test_text, ngram_range=(1,2),
5 #                                     min_df=2, keep_topN=500,
6 #                                     stop_words = stop_words,char_ngrams=False)
7
8 # #train
9 # train_vocab, train_df, train_count = get_vocab(train_text, ngram_range=(1,2),
10 #                                     min_df=10, keep_topN=100,
11 #                                     stop_words=stop_words,char_ngrams=False)
12 # #Dev
13 # dev_vocab, dev_df, dev_count = get_vocab(dev_text, ngram_range=(1,2),
14 #                                     min_df=10, keep_topN=100,
15 #                                     stop_words=stop_words,char_ngrams=False)
16
17 # #Test Vectorisation
18 # print("Vectorise test text....","\n")
19 # test_count, test_vect = vectorise(test_text, test_vocab)
20
21 # #Train Vectorisation
22 # print("Vectorise train text....","\n")
23 # train_count, train_vect = vectorise(train_text, train_vocab)
24
25 # #Dev Vectorisation
26 # print("Vectorise dev text....","\n")
27 # dev_count, dev_vect = vectorise(dev_text, dev_vocab)
28
29
30 weights_tfidf, training_loss_history_tfidf, validation_loss_history_tfidf = train_model(train_count, train_vect, dev_count, dev_vect, alpha=0.01, tolerance=0.001)
31
32
33
34 X_te_count = train_vect
35
36 w_count = weights_tfidf
37
38 preds_te_count = predict_class(X_te_count, w_count)
39
40 # train_count, weights
41
42
43 Y_te = dev_vect
44
45 print('Accuracy:', accuracy_score(Y_te,preds_te_count))
46 print('Precision:', precision_score(Y_te,preds_te_count))
```

```

47 print('Recall:', recall_score(Y_te,preds_te_count))
48 print('F1-Score:', f1_score(Y_te,preds_te_count))
49
50 # training_loss_history
51 # validation_loss_history
52
53 plt.figure(figsize=(25,6))
54
55 plt.title('Cost Function Slope')
56 plt.plot(training_loss_history_tfidf, label='Training Loss History')
57 plt.plot(validation_loss_history_tfidf, label='Validation Loss Hist')
58 plt.legend(prop={'size': 16})
59 plt.xlabel('Number of Iterations')
60 plt.ylabel('Error Values')
61 plt.show()
62
63 # print("Test Dictionary >>> ", "\n")
64 # id_w_test, w_id_test = create_2dict(test_df)
65
66 # print("Train Dictionary >>> ", "\n")
67 # id_w_train, w_id_train =create_2dict(train_df)
68
69 # print("DEV Dictionary >>> ", "\n")
70 # id_w_dev, w_id_dev =create_2dict(dev_df)
71
72
73
74 top_neg = w_count.argsort()[ :10]
75 for i in top_neg:
76     # print(id2word[i])
77     print(id_w_train[i])
78
79 top_pos = w_count.argsort()[::-1][ :10]
80 for i in top_pos:
81     # print(id2word[i])
82     print(id_w_train[i])

```

BOCN-count:

```

In [ ]: 1 ##### BOCN-count #####
2
3 # #TEST
4 # test_vocab, test_df, test_count = get_vocab(test_text, ngram_range=(1,2),
5 #                                           min_df=2, keep_topN=500,
6 #                                           stop_words = stop_words,char_ngrams=True)
7
8 #train
9 train_vocab_BOEN, train_df_BOEN, train_count_BOEN = get_vocab(train_text, ngram_range=(1,2),
10                                                                min_df=10, keep_topN=100,
11                                                                stop_words=stop_words,char_ngrams=True)
12 #Dev
13 dev_vocab_BOEN, dev_df_BOEN, dev_count_BOEN = get_vocab(dev_text, ngram_range=(1,2),
14                                                           min_df=10, keep_topN=100,
15                                                           stop_words=stop_words,char_ngrams=True)
16
17 # #Test Vectorisation
18 # print("Vectorise test text....","\n")
19 # test_count, test_vect = vectorise(test_text, test_vocab)
20

```

```

21 #Train Vectorisation
22 print("Vectorise train text....","\n")
23 train_count_BOCN, train_vect_BOCN = vectorise(train_text, train_voc
24
25 #Dev Vectorisation
26 print("Vectorise dev text....","\n")
27 dev_count_BOCN, dev_vect_BOCN = vectorise(dev_text, dev_vocab_BOCN)
28
29
30 weights_BOCN, training_loss_history_BOCN, validation_loss_history_F
31
32
33
34 X_te_count = train_count_BOCN
35
36 w_count = weights_BOCN
37
38 preds_te_count = predict_class(X_te_count, w_count)
39
40 # train_count, weights
41
42
43 Y_te = dev_count_BOCN
44
45 print('Accuracy:', accuracy_score(Y_te,preds_te_count))
46 print('Precision:', precision_score(Y_te,preds_te_count))
47 print('Recall:', recall_score(Y_te,preds_te_count))
48 print('F1-Score:', f1_score(Y_te,preds_te_count))
49
50 # training_loss_history
51 # validation_loss_history
52
53 # plt.figure(figsize=(25,6))
54
55 plt.title('Cost Function Slope')
56 plt.plot(training_loss_history_BOCN, label='Training Loss History')
57 plt.plot(validation_loss_history_BOCN, label='Validation Loss Histo
58 plt.legend(prop={'size': 16})
59 plt.xlabel('Number of Iterations')
60 plt.ylabel('Error Values')
61 plt.show()
62
63 # print("Test Dictionary >>> ", "\n")
64 # id_w_test, w_id_test = create_2dict(test_df)
65
66 print("Train Dictionary >>> ", "\n")
67 id_w_train_BOCN, w_id_train_BOCN =create_2dict(train_df_BOCN)
68
69 print("DEV Dictionary >>> ", "\n")
70 id_w_dev_BOCN, w_id_dev_BOCN =create_2dict(dev_df_BOCN)
71
72
73
74 top_neg = w_count.argsort()[::-10]
75 for i in top_neg:
76     print(id2word[i])
77     print(id_w_train_BOCN[i])
78
79 top_pos = w_count.argsort()[::-1][::-10]
80 for i in top_pos:

```

```

81 #     print(id2word[i])
82     print(id_w_train_BOEN[i])

```

BOEN-tfidf:

In []:

```

1 ##### BOCN-tfidf #####
2
3 # #TEST
4 # test_vocab, test_df, test_count = get_vocab(test_text, ngram_range=(1,2),
5 #                                           min_df=2, keep_topN=500,
6 #                                           stop_words = stop_words,char_ngrams=True)
7
8 #train
9 # train_vocab, train_df, train_count = get_vocab(train_text, ngram_range=(1,2),
10 #                                           min_df=10, keep_topN=100,
11 #                                           stop_words=stop_words,char_ngrams=True)
12 # #Dev
13 # dev_vocab, dev_df, dev_count = get_vocab(dev_text, ngram_range=(1,2),
14 #                                           min_df=10, keep_topN=100,
15 #                                           stop_words=stop_words,char_ngrams=True)
16
17 # #Test Vectorisation
18 # print("Vectorise test text....","\n")
19 # test_count, test_vect = vectorise(test_text, test_vocab)
20
21 #Train Vectorisation
22 # print("Vectorise train text....","\n")
23 # train_count, train_vect = vectorise(train_text, train_vocab)
24
25 # #Dev Vectorisation
26 # print("Vectorise dev text....","\n")
27 # dev_count, dev_vect = vectorise(dev_text, dev_vocab)
28
29
30 weights_BOEN_tfidf, training_loss_history_BOEN_tfidf, validation_loss_history_BOEN_tfidf = train_model(train_count, train_vect, dev_count, dev_vect, test_count, test_vect)
31
32
33
34 X_te_count = train_vect_BOEN_tfidf
35
36 w_count = weights_BOEN_tfidf
37
38 preds_te_count = predict_class(X_te_count, w_count)
39
40 # train_count, weights
41
42
43 Y_te = dev_vect_BOEN
44
45 print('Accuracy:', accuracy_score(Y_te,preds_te_count))
46 print('Precision:', precision_score(Y_te,preds_te_count))
47 print('Recall:', recall_score(Y_te,preds_te_count))
48 print('F1-Score:', f1_score(Y_te,preds_te_count))
49
50 # training_loss_history
51 # validation_loss_history
52
53 plt.figure(figsize=(25,6))
54

```

```

55 plt.title('Cost Function Slope')
56 plt.plot(training_loss_history_BOCN_tfidf, label='Training Loss His
57 plt.plot(validation_loss_history_BOCN_tfidf, label='Validation Loss
58 plt.legend(prop={'size': 16})
59 plt.xlabel('Number of Iterations')
60 plt.ylabel('Error Values')
61 plt.show()
62
63 # print("Test Dictionary >>> ", "\n")
64 # id_w_test, w_id_test = create_2dict(test_df)
65
66 # print("Train Dictionary >>> ", "\n")
67 # id_w_train, w_id_train =create_2dict(train_df)
68
69 # print("DEV Dictionary >>> ", "\n")
70 # id_w_dev, w_id_dev =create_2dict(dev_df)
71
72
73
74 top_neg = w_count.argsort()[ :10]
75 for i in top_neg:
76     # print(id2word[i])
77     print(id_w_train_BOCN_tfidf[i])
78
79 top_pos = w_count.argsort()[::-1][ :10]
80 for i in top_pos:
81     # print(id2word[i])
82     print(id_w_train_BOCN_tfidf[i])

```

BOW+BOCN:

In []:

```
1 # BOW+BOCN
```

Full Results

Add here your results:

LR	Precision	Recall	F1-Score
BOW-count			
BOW-tfidf			
BOCN-count			
BOCN-tfidf			
BOW+BOCN			

Please discuss why your best performing model is better than the rest.

In []:

```
1 # BOW+BOCN
```

In []:

```
1
```

